

House Price Prediction

Importing Libraries and Data Loading

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df=pd.read_csv("/content/House price Prediction.csv")
df
```

```
Out[2]:
```

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view
0	2014-05-02 00:00:00	3.130000e+05	3.0	1.50	1340	7912	1.5	0	C
1	2014-05-02 00:00:00	2.384000e+06	5.0	2.50	3650	9050	2.0	0	4
2	2014-05-02 00:00:00	3.420000e+05	3.0	2.00	1930	11947	1.0	0	C
3	2014-05-02 00:00:00	4.200000e+05	3.0	2.25	2000	8030	1.0	0	C
4	2014-05-02 00:00:00	5.500000e+05	4.0	2.50	1940	10500	1.0	0	C
...
4595	2014-07-09 00:00:00	3.081667e+05	3.0	1.75	1510	6360	1.0	0	C
4596	2014-07-09 00:00:00	5.343333e+05	3.0	2.50	1460	7573	2.0	0	C
4597	2014-07-09 00:00:00	4.169042e+05	3.0	2.50	3010	7014	2.0	0	C
4598	2014-07-10 00:00:00	2.034000e+05	4.0	2.00	2090	6630	1.0	0	C
4599	2014-07-10 00:00:00	2.206000e+05	3.0	2.50	1490	8102	2.0	0	C

4600 rows × 18 columns

```
In [3]: df.head()
```

Out[3]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	cond
0	2014-05-02 00:00:00	313000.0	3.0	1.50	1340	7912	1.5	0	0	
1	2014-05-02 00:00:00	2384000.0	5.0	2.50	3650	9050	2.0	0	4	
2	2014-05-02 00:00:00	342000.0	3.0	2.00	1930	11947	1.0	0	0	
3	2014-05-02 00:00:00	420000.0	3.0	2.25	2000	8030	1.0	0	0	
4	2014-05-02 00:00:00	550000.0	4.0	2.50	1940	10500	1.0	0	0	

In [4]:

df.tail()

Out[4]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	cond
4595	2014-07-09 00:00:00	308166.666667	3.0	1.75	1510	6360	1.0	0		
4596	2014-07-09 00:00:00	534333.333333	3.0	2.50	1460	7573	2.0	0		
4597	2014-07-09 00:00:00	416904.166667	3.0	2.50	3010	7014	2.0	0		
4598	2014-07-10 00:00:00	203400.000000	4.0	2.00	2090	6630	1.0	0		
4599	2014-07-10 00:00:00	220600.000000	3.0	2.50	1490	8102	2.0	0		

In [5]:

df.isna().sum()

Out[5]:

	0
date	0
price	0
bedrooms	0
bathrooms	0
sqft_living	0
sqft_lot	0
floors	0
waterfront	0
view	0
condition	0
sqft_above	0
sqft_basement	0
yr_built	0
yr_renovated	0
street	0
city	0
statezip	0
country	0

dtype: int64

In [6]:

df.dtypes

Out[6]: 0

date	object
price	float64
bedrooms	float64
bathrooms	float64
sqft_living	int64
sqft_lot	int64
floors	float64
waterfront	int64
view	int64
condition	int64
sqft_above	int64
sqft_basement	int64
yr_built	int64
yr_renovated	int64
street	object
city	object
statezip	object
country	object

dtype: object

```
In [7]: df.drop(['date', 'street', 'city', 'country'], axis=1, inplace=True)
df
```

Out[7]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	conditi
0	3.130000e+05	3.0	1.50	1340	7912	1.5	0	0	
1	2.384000e+06	5.0	2.50	3650	9050	2.0	0	4	
2	3.420000e+05	3.0	2.00	1930	11947	1.0	0	0	
3	4.200000e+05	3.0	2.25	2000	8030	1.0	0	0	
4	5.500000e+05	4.0	2.50	1940	10500	1.0	0	0	
...
4595	3.081667e+05	3.0	1.75	1510	6360	1.0	0	0	
4596	5.343333e+05	3.0	2.50	1460	7573	2.0	0	0	
4597	4.169042e+05	3.0	2.50	3010	7014	2.0	0	0	
4598	2.034000e+05	4.0	2.00	2090	6630	1.0	0	0	
4599	2.206000e+05	3.0	2.50	1490	8102	2.0	0	0	

4600 rows × 14 columns



```
In [8]: #Showing no.of unique values in each columns
df.nunique()
```

Out[8]: 0

price	1741
bedrooms	10
bathrooms	26
sqft_living	566
sqft_lot	3113
floors	6
waterfront	2
view	5
condition	5
sqft_above	511
sqft_basement	207
yr_built	115
yr_renovated	60
statezip	77

dtype: int64

```
In [9]: for i in df['yr_renovated']:
        if i!=0:
            df['yr_renovated']=df['yr_renovated'].replace(i,2024-i)
df
```

Out[9]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	conditi
0	3.130000e+05	3.0	1.50	1340	7912	1.5	0	0	
1	2.384000e+06	5.0	2.50	3650	9050	2.0	0	4	
2	3.420000e+05	3.0	2.00	1930	11947	1.0	0	0	
3	4.200000e+05	3.0	2.25	2000	8030	1.0	0	0	
4	5.500000e+05	4.0	2.50	1940	10500	1.0	0	0	
...
4595	3.081667e+05	3.0	1.75	1510	6360	1.0	0	0	
4596	5.343333e+05	3.0	2.50	1460	7573	2.0	0	0	
4597	4.169042e+05	3.0	2.50	3010	7014	2.0	0	0	
4598	2.034000e+05	4.0	2.00	2090	6630	1.0	0	0	
4599	2.206000e+05	3.0	2.50	1490	8102	2.0	0	0	

4600 rows × 14 columns

```
In [10]: df['statezip']=df['statezip'].str.replace('WA','')
df['statezip']=df['statezip'].astype(int)
df
```

Out[10]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition
0	3.130000e+05	3.0	1.50	1340	7912	1.5	0	0	
1	2.384000e+06	5.0	2.50	3650	9050	2.0	0	4	
2	3.420000e+05	3.0	2.00	1930	11947	1.0	0	0	
3	4.200000e+05	3.0	2.25	2000	8030	1.0	0	0	
4	5.500000e+05	4.0	2.50	1940	10500	1.0	0	0	
...
4595	3.081667e+05	3.0	1.75	1510	6360	1.0	0	0	
4596	5.343333e+05	3.0	2.50	1460	7573	2.0	0	0	
4597	4.169042e+05	3.0	2.50	3010	7014	2.0	0	0	
4598	2.034000e+05	4.0	2.00	2090	6630	1.0	0	0	
4599	2.206000e+05	3.0	2.50	1490	8102	2.0	0	0	

4600 rows × 14 columns

In [11]: df.dtypes

Out[11]:


	0
price	float64
bedrooms	float64
bathrooms	float64
sqft_living	int64
sqft_lot	int64
floors	float64
waterfront	int64
view	int64
condition	int64
sqft_above	int64
sqft_basement	int64
yr_built	int64
yr_renovated	int64
statezip	int64

dtype: object

```
In [12]: df1 = df.copy()
#Correlation with Response Variable class
X = df1.drop(['price'],axis=1)
y = df1['price']
X.corrwith(y).plot.bar()
```



```
figsize = (16, 5), title = "Correlation with Response Variable class", fontsize = 14,
rot = 45, grid = False)
plt.show()
```




```
In [13]: df.corr()
```

Out[13]:



	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	sqft_above	sqft_basement	yr_built	yr_renovated	statezip
price	1.000000	0.200336	0.327110	0.430410	0.050451	0.151461	0.135648	0.228504	0.034915	0.367570	0.210427	0.021857	-0.011928	-0.046052
bedrooms	0.200336	1.000000	0.545920	0.594884	0.068819	0.177895	-0.003483	0.111028	0.025080	0.484705	0.334165	0.142461	-0.026228	-0.153443
bathrooms	0.327110	0.545920	1.000000	0.761154	0.107837	0.486428	0.076232	0.211960	-0.119994	0.689918	0.298020	0.463498	-0.108796	-0.206231
sqft_living	0.430410	0.594884	0.761154	1.000000	0.210538	0.344850	0.117616	0.311009	-0.062826	0.876443	0.447206	0.287775	-0.061498	-0.210891
sqft_lot	0.050451	0.068819	0.107837	0.210538	1.000000	0.003750	0.017241	0.073907	0.000558	0.216455	0.034842	0.050706	-0.039293	-0.133509
floors	0.151461	0.177895	0.486428	0.344850	0.003750	1.000000	0.022024	0.031211	-0.275013	0.522814	-0.255510	0.467481	-0.089938	-0.064999
waterfront	0.135648	-0.003483	0.076232	0.117616	0.017241	0.022024	1.000000	0.360935	0.000352	0.078911	0.097501	-0.023563	0.000165	0.005937
view	0.228504	0.111028	0.211960	0.311009	0.073907	0.031211	0.360935	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
condition	0.034915	0.025080	-0.119994	-0.062826	0.000558	-0.275013	0.000352	0.000000	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000
sqft_above	0.367570	0.484705	0.689918	0.876443	0.216455	0.522814	0.078911	0.000000	0.000000	1.000000	0.000000	0.000000	0.000000	0.000000
sqft_basement	0.210427	0.334165	0.298020	0.447206	0.034842	-0.255510	0.097501	0.000000	0.000000	0.000000	1.000000	0.000000	0.000000	0.000000
yr_built	0.021857	0.142461	0.463498	0.287775	0.050706	0.467481	-0.023563	0.000000	0.000000	0.000000	0.000000	1.000000	0.000000	0.000000
yr_renovated	-0.011928	-0.026228	-0.108796	-0.061498	-0.039293	-0.089938	0.000165	0.000000	0.000000	0.000000	0.000000	0.000000	1.000000	0.000000
statezip	-0.046052	-0.153443	-0.206231	-0.210891	-0.133509	-0.064999	0.005937	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	1.000000



```
In [14]: plt.figure(figsize=(10,10))
sns.heatmap(df.corr(),annot=True)
plt.show()
```







Removal of Outliers



```
In [15]: for i in df.columns:
plt.figure(figsize=(6,6))
sns.boxplot(df[i],color='blue')
plt.xlabel(i)
plt.ylabel('Range')
```



















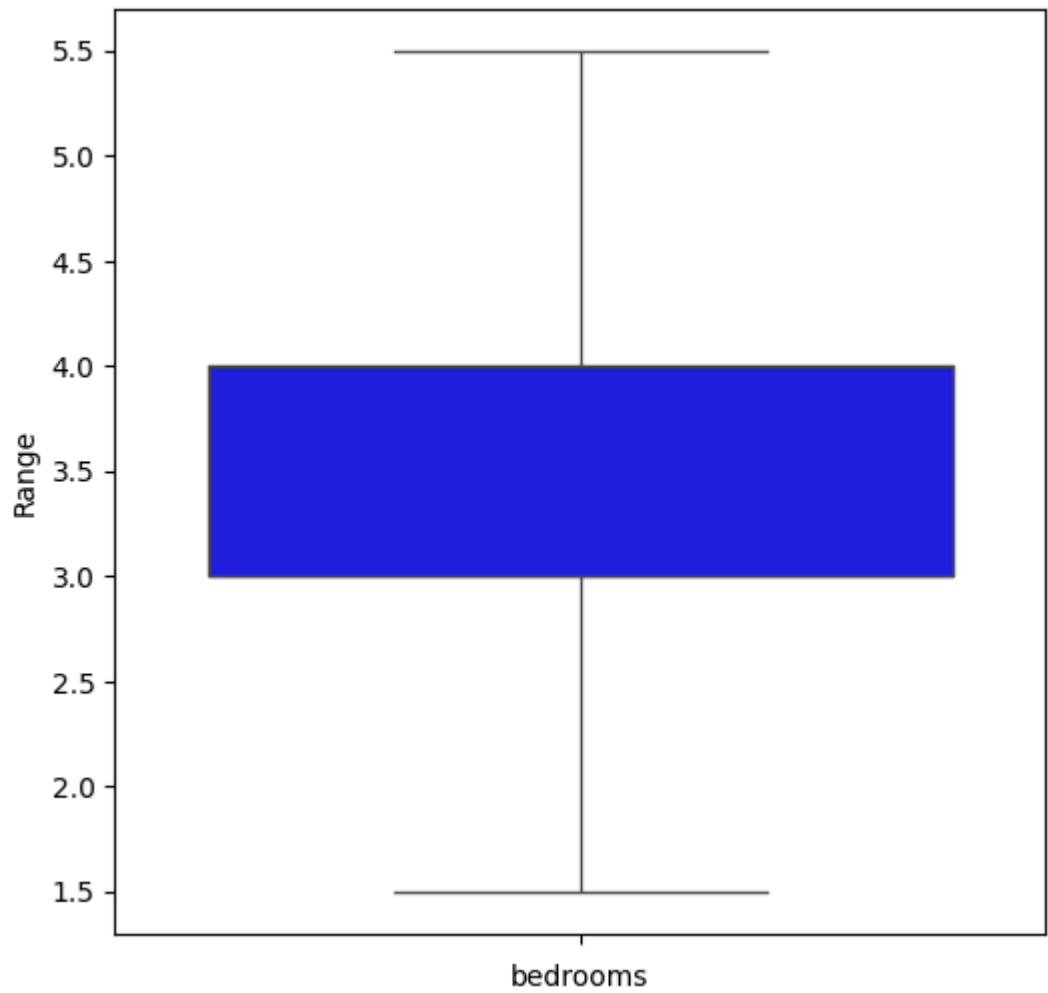




```
In [16]: df.drop(['waterfront', 'view'], axis=1, inplace=True)
```



```
In [17]: for i in df.columns:
def iqr_method(df, variables):
    q1=df[variables].quantile(0.25)
    q3=df[variables].quantile(0.75)
    iqr=q3-q1
    upper=q3+(1.5*iqr)
    lower=q1-(1.5*iqr)
    return lower, upper
lower_lim, upper_lim=iqr_method(df, i)
df[i]=np.where(df[i]>upper_lim, upper_lim, np.where(df[i]<lower_lim, lower_lim, df[i]))
```



```
In [18]: for i in df.columns:
plt.figure(figsize=(6,6))
sns.boxplot(df[i], color='blue')
plt.xlabel(i)
plt.ylabel('Range')
```


















```
In [19]: x=df.drop(['price'],axis=1)
y=df['price']
x
```

```
Out[19]:
```

	bedrooms	bathrooms	sqft_living	sqft_lot	floors	condition	sqft_above	sqft_basement	y
0	3.0	1.50	1340.0	7912.0	1.5	3.0	1340.0	0.0	
1	5.0	2.50	3650.0	9050.0	2.0	5.0	3370.0	280.0	
2	3.0	2.00	1930.0	11947.0	1.0	4.0	1930.0	0.0	
3	3.0	2.25	2000.0	8030.0	1.0	4.0	1000.0	1000.0	
4	4.0	2.50	1940.0	10500.0	1.0	4.0	1140.0	800.0	
...
4595	3.0	1.75	1510.0	6360.0	1.0	4.0	1510.0	0.0	
4596	3.0	2.50	1460.0	7573.0	2.0	3.0	1460.0	0.0	
4597	3.0	2.50	3010.0	7014.0	2.0	3.0	3010.0	0.0	
4598	4.0	2.00	2090.0	6630.0	1.0	3.0	1070.0	1020.0	
4599	3.0	2.50	1490.0	8102.0	2.0	4.0	1490.0	0.0	

4600 rows × 11 columns

```
In [20]: y
```



Out[20]:



	price
0	3.130000e+05
1	1.153094e+06
2	3.420000e+05
3	4.200000e+05
4	5.500000e+05
...	...
4595	3.081667e+05
4596	5.343333e+05
4597	4.169042e+05
4598	2.034000e+05
4599	2.206000e+05



4600 rows × 1 columns



dtype: float64



```
In [21]: for i in x.columns:  
          plt.scatter(x=y,y=x[i])  
          plt.xlabel('price')  
          plt.ylabel(i)  
          plt.show()
```















```
In [22]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.30,random_state=1)
x_train
```

Out[22]:

	bedrooms	bathrooms	sqft_living	sqft_lot	floors	condition	sqft_above	sqft_basement	y
3990	4.0	3.250	3990.0	9786.0	2.0	3.0	3965.0	0.0	
1507	4.0	2.500	2680.0	7178.0	2.0	3.0	2680.0	0.0	
1652	3.0	2.250	2675.0	20002.0	2.0	3.0	2675.0	0.0	
2279	5.5	3.625	3440.0	4500.0	2.0	3.0	3280.0	160.0	
2106	4.0	1.750	2220.0	6500.0	2.0	4.0	2220.0	0.0	
...
2895	4.0	1.500	1220.0	9600.0	1.0	3.0	1220.0	0.0	
2763	3.0	1.750	1970.0	8200.0	1.0	5.0	1420.0	550.0	
905	3.0	2.500	1720.0	1916.0	2.0	3.0	1720.0	0.0	
3980	4.0	3.250	4100.0	20002.0	2.0	3.0	2500.0	1525.0	
235	3.0	2.750	1540.0	6760.0	1.0	5.0	1210.0	330.0	

3220 rows × 11 columns

In [23]: x_test

Out[23]:

	bedrooms	bathrooms	sqft_living	sqft_lot	floors	condition	sqft_above	sqft_basement	y
1351	5.0	2.000	2330.0	10750.0	1.0	4.0	1190.0	1140.0	
1687	3.0	2.500	2730.0	5832.0	2.0	3.0	2730.0	0.0	
1297	3.0	2.000	1220.0	1186.0	3.0	3.0	1220.0	0.0	
2101	4.0	1.750	2700.0	7875.0	1.5	4.0	2700.0	0.0	
3920	3.0	3.625	2080.0	2250.0	3.0	3.0	2080.0	0.0	
...
3490	4.0	2.500	2250.0	7526.0	2.0	3.0	2250.0	0.0	
3584	2.0	0.750	840.0	20002.0	1.0	4.0	840.0	0.0	
855	2.0	2.000	1360.0	4688.0	1.0	3.0	780.0	580.0	
309	5.0	3.625	4270.0	8076.0	2.0	3.0	3400.0	870.0	
1929	3.0	2.500	2680.0	9750.0	1.0	4.0	1610.0	1070.0	

1380 rows × 11 columns

In [24]: y_train

Out[24]:

	price
3990	866000.0
1507	515000.0
1652	660000.0
2279	495000.0
2106	480000.0
...	...
2895	310000.0
2763	540000.0
905	280000.0
3980	775000.0
235	503000.0

3220 rows × 1 columns

dtype: float64

In [25]: y_test

Out[25]:

	price
1351	389000.00
1687	1000000.00
1297	355000.00
2101	390000.00
3920	715000.00
...	...
3490	440000.00
3584	528000.00
855	488000.00
309	1153093.75
1929	653000.00

1380 rows × 1 columns

dtype: float64*Model Creation and Performance Evaluation*

```
In [26]: from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import r2_score
linreg=LinearRegression()
```



```
ranfor=RandomForestRegressor()  
dectree=DecisionTreeRegressor()  
grad=GradientBoostingRegressor()  
lst=[linreg,ranfor,dectree,grad]
```

```
In [27]: for i in lst:  
         i.fit(x_train,y_train)  
         y_pred=i.predict(x_test)  
         print("R2score of",i,"is",r2_score(y_test,y_pred))
```

```
R2score of LinearRegression() is 0.483250199383935  
R2score of RandomForestRegressor() is 0.6058197511228784  
R2score of DecisionTreeRegressor() is 0.33683930554751207  
R2score of GradientBoostingRegressor() is 0.6351632169487991
```